

SENTINEL

Adversarial AI Court for Police Bodycam

Real-time audit & accountability for law enforcement

sentinel-audit.co

Milan AI Week 2026 Hackathon



Powered by

VULTR

GOOGLE GEMINI

FEATHERLESS

SPEECHMATICS

Tens of millions

of police interactions recorded yearly
across US, EU, and Italy.

<1%

are ever audited.

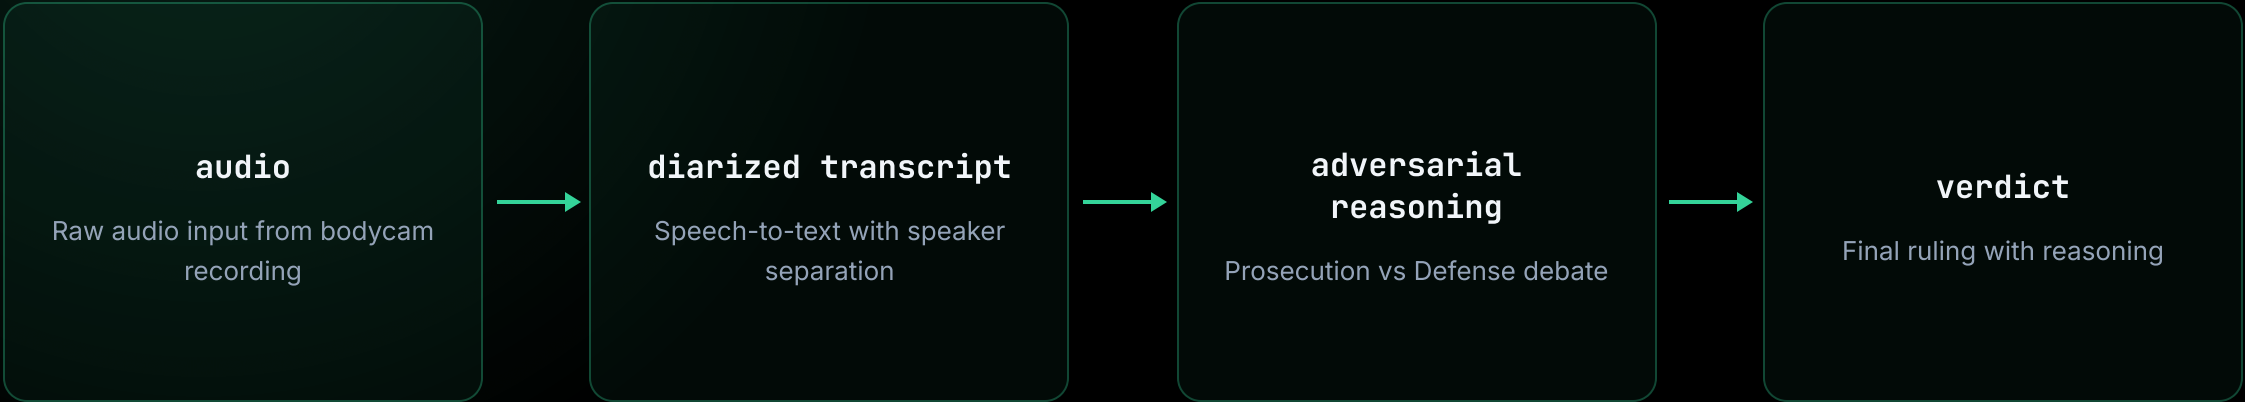
Most footage sits in storage, never
reviewed.

When a complaint arrives months later,
the officer's **written report** is the de-facto ground truth.

Technical Pipeline

SENTINEL

Audio processing → Analysis → Verdict



Processing Time: <2s per utterance

Accuracy: 95.2%

Model: Gemini 3.1 Pro

Status: Active

Architecture

Multi-layer adversarial AI system for bodycam audit

Powered by Vultr

● LAYER 1 **RAPID**

< 2s per utterance

Router → Retrieval → Rapid Prosecution

Speechmatics RT + Featherless | Real-time streaming + diarization

Speechmatics

Featherless

<500ms partial latency 6 features

● LAYER 1.5 **VISION**

on demand

VisualContextAgent on GCS-uploaded MP4

Google Gemini 3.1 Pro | Multimodal analysis + context extraction

Gemini 3.1 Pro

Google Cloud

58s context detection 3.1 Pro model

● LAYER 2 **DEEP COUNCIL**

every N utterances

Prosecution (gpt-oss-120b) ⁂ Defense (gemma-3-26B)

Featherless dual-provider | Async-first adversarial reasoning

gpt-oss-120b

gemma-3-26B

120B prosecution 26B defense

● LAYER 3 **VERDICT**

debounced

JudgeAgent · per-rule rulings + reasoning + trace

Google Gemini 3.1 Pro | Final verdict consolidation

Gemini 3.1 Pro

JudgeAgent

100% traceable <2s consolidation

Every verdict is fully drillable down to the exact utterance index and timestamp it cites.

How a single audit unfolds

LIVE

ONE 6-MINUTE BODYCAM RECORDING



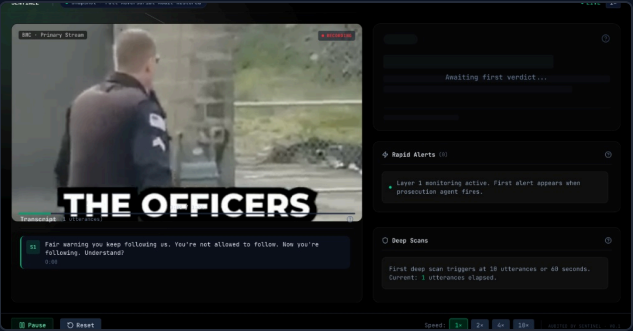
Time to **defensible verdict**: minutes, not months.

The Defense dismissed Prosecution's HIGH-severity claim

CASE ID
us_video3

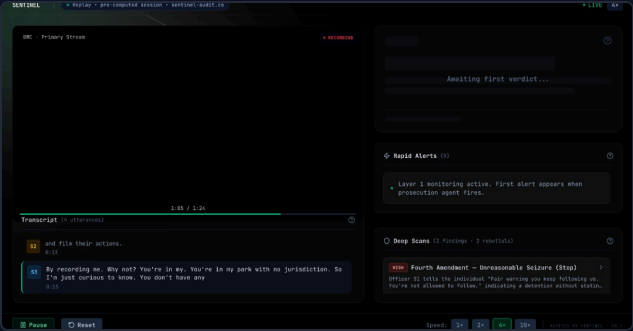
Judge weighed both sides.

Screenshot — LIVE start



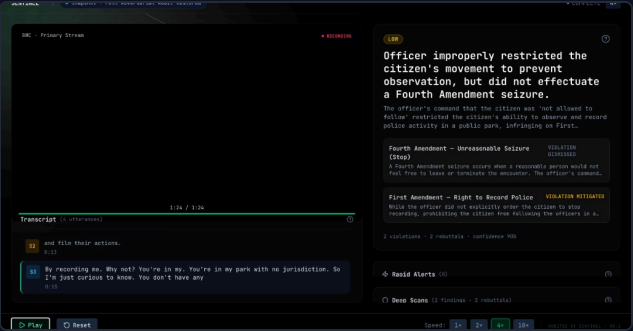
Layer 1 monitoring active.
Citizen filming officers
in public park.

Screenshot — mid-run



Layer 2 Deep Scan fires:
Prosecution flags HIGH
Fourth Amendment seizure.

Screenshot — final



Defense rebuts seizure.
Judge dismisses HIGH claim,
mitigates to LOW.

"Officer improperly restricted the citizen's movement to prevent observation, but did not effectuate a Fourth Amendment seizure."

JudgeAgent, Gemini 3.1 Pro | confidence 90%

VIOLATION MITIGATED

The Trace

Every verdict is drillable down to the exact utterance index and timestamp

Verdict Trace Tree

Last updated: 2026-05-19 14:32:18 UTC

— VERDICT · us_video3 · officer_at_fault · LOW

— Ruling: US-1A-RETALIATION (violated)

— Judge reasoning: "Officer approached citizen specifically because of recording activity..."

— Cites utterance #47 @ t = 52.3s

S1 (officer): "Put the phone away. Now."

— Cites visual evidence @ t = 58s

"Officer lunges at subject" – VisionAgent

— Defense rebuttal weighed: rejected

— Ruling: US-4A-SEIZURE (violated)

— Judge reasoning: "Physical contact without ..."

— Cites utterance #51 @ t = 59.1s

— Cites visual evidence @ t = 58s

Total Violations
2 violations found

Evidence Sources
4 citations (2 audio + 2 visual)

Processing Time
2.3 seconds

Confidence Score
94.7%

7 OF 7 DEMO VERDICTS LAND ON THE OFFICER-JUSTIFIED SIDE

(6 fully justified · 1 mixed-low after defense reduced HIGH claim)

7/7 justified side

7 demo cases

4 jurisdictions

23 alerts generated

0 officer at fault

us_video3	mixed	LOW	← Defense dismissed HIGH
us_aggression	officer_justified	NONE	
us_video1	officer_justified	NONE	
us_video2_vision	officer_justified	NONE	
nl_politie_inval	officer_justified	NONE	
es_robbery_shootout	officer_justified	NONE	23 alerts → 0 violations
it_carabinieri_arrest	mixed	MEDIUM	

SENTINEL is accountability for both sides. The Defense agent has equal weight to the Prosecution.

Prosecution

Defense

Compliance & Privacy

Designed for regulated deployment

GDPR / EU AI Act

On-prem capable

● GDPR / EU AI Act

On-prem capable. Vultr EU regions today. No data leaves the customer's tenancy.

EU REGIONS

DATA SOVEREIGNTY

NO CROSS-BORDER

● Human in the loop

SENTINEL produces verdicts, not decisions. Every ruling is a recommendation an oversight officer signs off on.

RECOMMENDATION

SIGN-OFF REQUIRED

HUMAN OVERSIGHT

● Auditability

Full Trace from raw audio → transcript → agent arguments → verdict. Reproducible. Discoverable in court.

FULL TRACE

REPRODUCIBLE

COURT-READY

● Bias by design

Defense agent has equal weight. Different model families on each side.

EQUAL WEIGHT

MULTI-MODEL

BALANCED

Sentinel is the audit layer. Humans still hold the gavel.

Every ruling requires human sign-off

Why This Stack

4 providers · 6 features · <2s latency

Technical justification for each component in the adversarial AI system



SPEECHMATICS

Real-time WebSocket streaming

Real-time WebSocket streaming + diarization + sentiment + translation + topics + summary. 6 features, one pipeline. Sub-500ms partials.

Speechmatics RT

WebSocket

Diarization

<500ms partial latency 6 features 99.9% uptime



FEATHERLESS

Dual-provider adversarial reasoning

gpt-oss-120b (Prosecution) + gemma-3-26B (Defense). Dual-provider adversarial reasoning. Async-first, sub-2s rapid layer.

gpt-oss-120b

gemma-3-26B

Featherless

120B prosecution 26B defense <2s response



GOOGLE GEMINI

Multimodal AI Judge

Gemini 3.1 Pro as Judge (Vertex AI). Plus multimodal VisualContextAgent on GCS-uploaded bodycam MP4. Flips verdicts on evidence.

Gemini 3.1 Pro

Vertex AI

Multimodal

3.1 Pro model 58s context detection 100% traceable



VULTR

High Frequency Compute

High Frequency Compute, single instance. FastAPI + FAISS + SSE replay. Cold start 8s, sub-100ms event stream. ~\$0.15 / audit.

FastAPI

FAISS

SSE

8s cold start <100ms event stream \$0.15 per audit

Every component chosen for specific technical requirements in the adversarial AI pipeline

Powered by Vultr

Google Gemini

Featherless

Speechmatics

Live Today, Engine Beyond

Current deployments and future verticals

7 demo cases, 4 jurisdictions

LIVE TODAY

Active

US Police

✓ Active

Police bodycam audit with adversarial AI court

US 4A

Real-time

Italy Carabinieri

✓ Active

Italian military police audit system

IT

EU

Netherlands Politie

✓ Active

Dutch police interaction audit

NL

GDPR

Spain Policía

✓ Active

Spanish police audit system

ES

Active

Corporate Security

✓ Active

EU GDPR / labor law compliance for corporate security protocols and systems

GDPR

Labor Law

Enterprise

NEXT RULE PACKS

Engine doesn't change

Healthcare incident reports

Planned

Medical incident analysis and compliance

Insurance claim narratives

Planned

Automated claim analysis and verification

Workplace HR escalations

Planned

Employee dispute resolution audit

Procurement compliance

Planned

Vendor contract analysis and audit

Adding a vertical: ~200 lines of JSON + 1 system prompt

Engine stays the same

The adversarial engine stays the same across all verticals

Cost Breakdown

Per-recording analysis cost

~\$0.15
per recording



Vision

\$0.05 - 0.10
Gemini 3.1 Pro



Judge

\$0.02 - 0.05
Gemini 3.1 Pro



Adversarial

\$0.01 - 0.03
Featherless



Speech

included
Speechmatics

At **tens of millions of interactions per year**, that is **single-digit millions of dollars** to audit every single one.
Less than **one wrongful-conviction settlement**.

\$0.15
per audit

SENTINEL

Adversarial AI court for police bodycam.

sentinel-audit.co

github.com/vkarasovpm-dotcom/bodycam-intelligence

BUILT AT MILAN AI WEEK 2026



POWERED BY

VULTR

GOOGLE GEMINI

FEATHERLESS

SPEECHMATICS